

# electricity\_pricing\_model

August 30, 2022

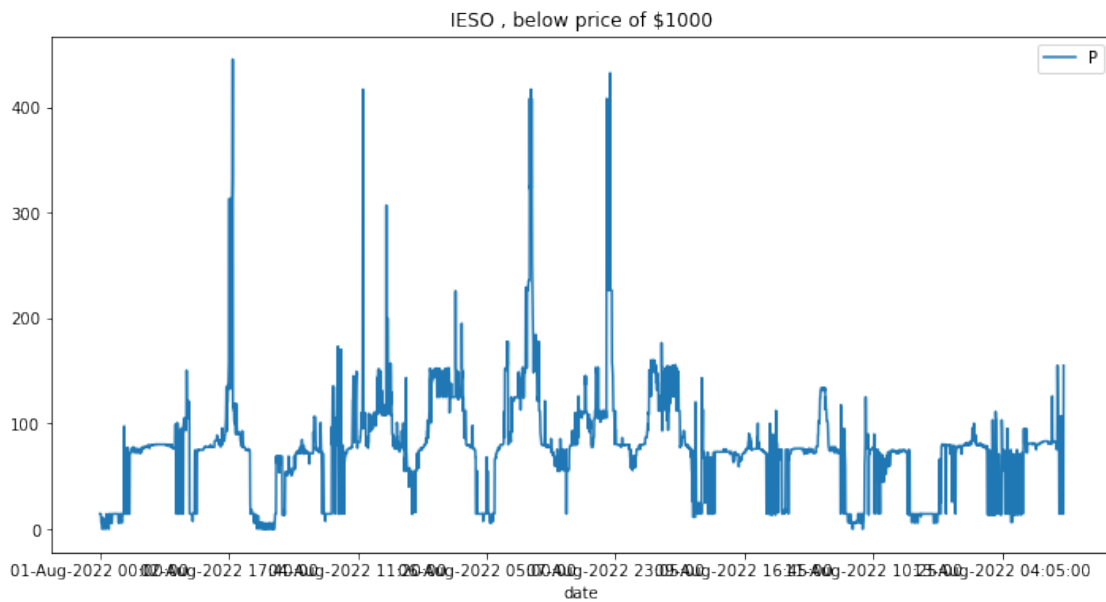
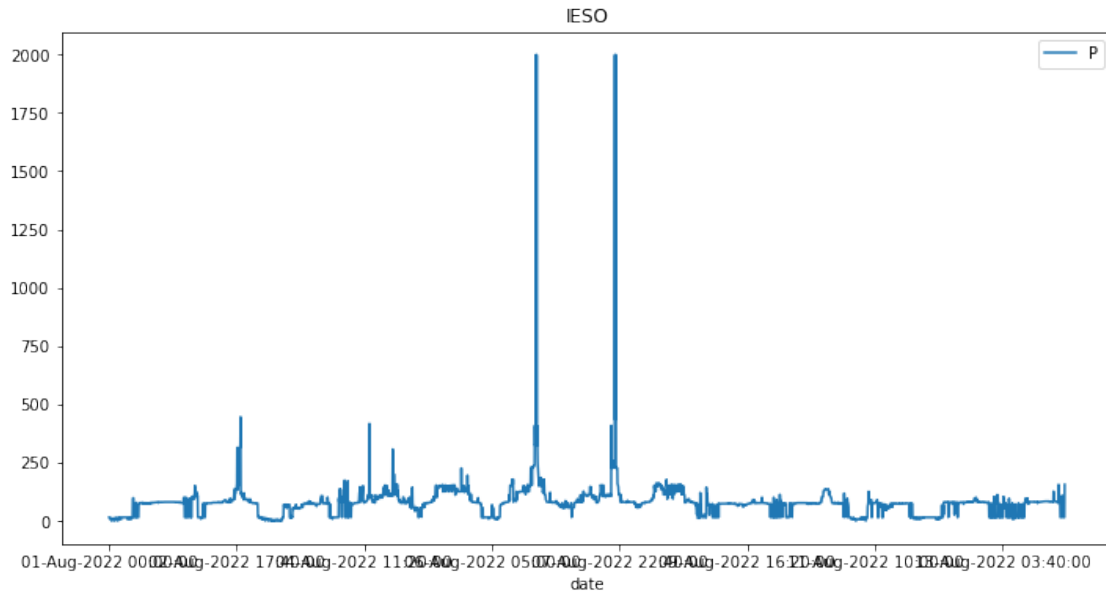
```
[2]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
from sklearn.linear_model import Ridge
from sklearn.model_selection import cross_val_score

[3]: df_ieso = pd.read_csv("dataset_ieso.csv") # aug 1 to 13, 3745 x 4
df_nyiso = pd.read_csv("dataset_nyiso.csv") # jan 1 to jan 31, 9300 x 4
df_pjm = pd.read_csv("dataset_pjm.csv") # may to may 13, 4032 x 5
df_nyiso2 = pd.read_csv("dataset_nyiso2.csv") # jan 1 to jan 31, 9300 x 4, ␣
    ↪ relabelled column names
#print(df_pjm)
#df_ieso.head()
#df_ieso.shape

[4]: #plot IESO
fig, ax = plt.subplots(figsize=(12, 6))
df_ieso.plot('date', 'P', ax=ax)
ax.set(title="IESO")

fig, ax = plt.subplots(figsize=(12, 6)) # plots with the two big outliers ␣
    ↪ removed
df_ieso[df_ieso["P"] < 1000].plot('date', 'P', ax=ax)
ax.set(title="IESO , below price of $1000")

[4]: [Text(0.5, 1.0, 'IESO , below price of $1000')]
```



```
[5]: #plot NYISO
fig, ax = plt.subplots(figsize=(12, 6))
df_nyiso.plot('date', 'rt_lmp', ax=ax)
ax.set(title="NYISO")

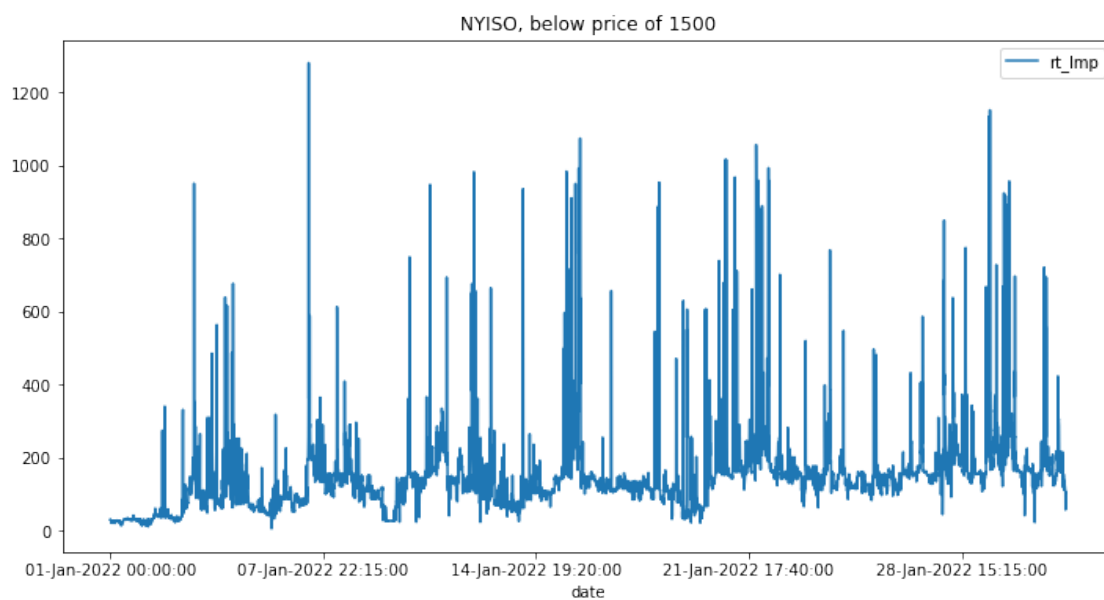
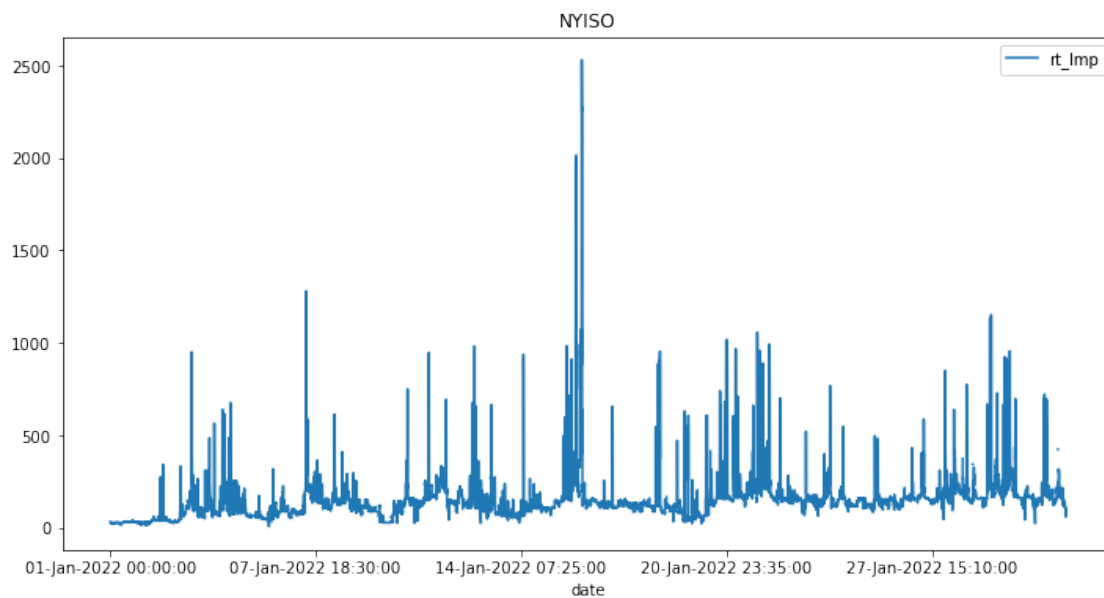
fig, ax = plt.subplots(figsize=(12, 6))
df_nyiso[df_nyiso["rt_lmp"] < 1500].plot('date', 'rt_lmp', ax=ax)
```

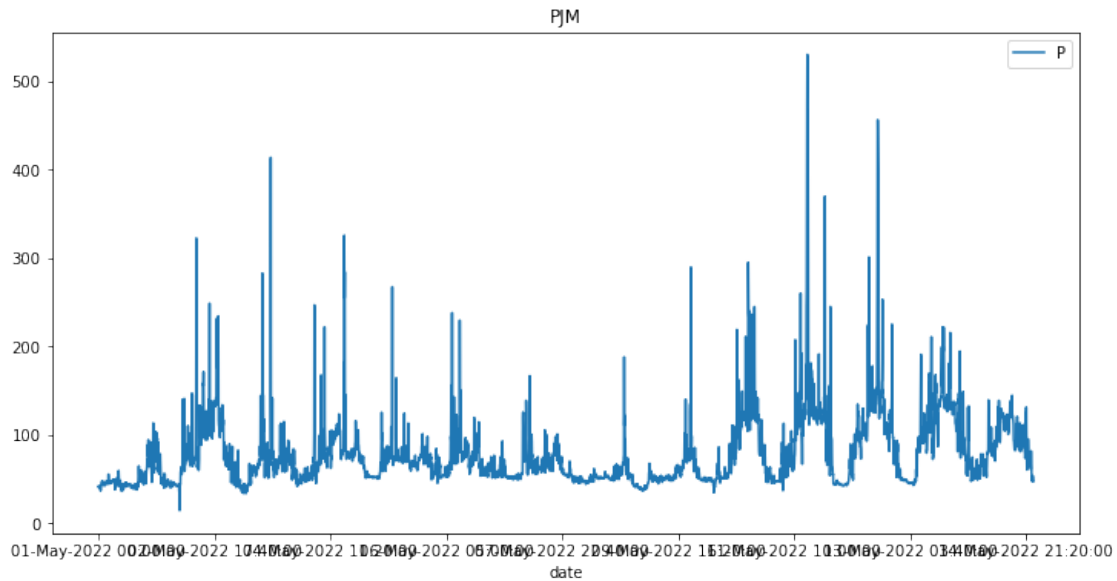
```

ax.set(title="NYISO, below price of 1500")
#plot PJM
fig, ax = plt.subplots(figsize=(12, 6))
df_pjm.plot('date', 'P', ax=ax)
ax.set(title="PJM")

```

[5]: [Text(0.5, 1.0, 'PJM')]





```
[6]: df_ieso.describe() #summary stats
```

```
[6]:
```

	P	G	L
count	3745.000000	3745.000000	3745.000000
mean	77.687554	7.035543	17231.581308
std	84.277547	0.333045	2628.716183
min	0.000000	6.340000	12250.000000
25%	57.900000	6.870000	15034.000000
50%	76.280000	7.100000	17282.000000
75%	86.590000	7.260000	19455.000000
max	1999.000000	7.630000	22062.000000

```
[7]: df_ieso.isna().sum() # is their missing data?
```

```
[7]: date    0
      P      0
      G      0
      L      0
      dtype: int64
```

```
[8]: df_nyiso.describe() #summary stats
```

```
[8]:
```

	rt_lmp	gas	load_rt
count	8989.000000	8989.000000	8983.000000
mean	146.399716	11.202847	18777.870979
std	123.396331	4.964135	2003.687307
min	6.080000	3.540000	13063.000000
25%	90.300000	7.150000	17302.000000

50%	134.330000	12.020000	19144.000000
75%	163.040000	14.460000	20225.500000
max	2529.890000	20.000000	23424.000000

```
[9]: df_nyiso.isna().sum() # is their missing data?
```

```
[9]: date      0
      rt_lmp    311
      gas      311
      load_rt   317
      dtype: int64
```

```
[10]: df_pjm.describe() #summary stats
```

```
[10]:
```

	P	G	L
count	4032.000000	4032.000000	4019.000000
mean	77.533418	7.006429	78318.279920
std	40.109895	0.608218	9333.851883
min	14.530000	6.020000	61403.000000
25%	51.640000	6.720000	70264.000000
50%	66.760000	6.890000	79059.000000
75%	91.575000	7.640000	84125.500000
max	529.820000	7.710000	101092.000000

```
[11]: df_pjm.isna().sum() # is their missing data?
```

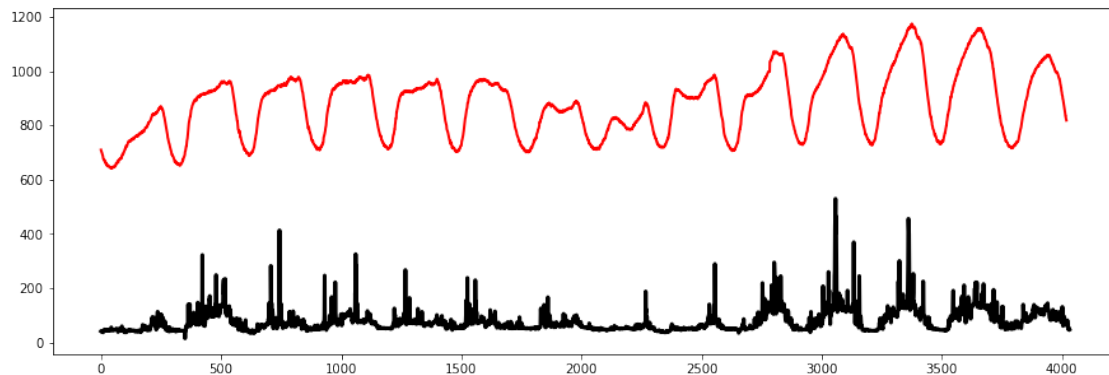
```
[11]: date      0
      P        0
      G        0
      L      13
      dtype: int64
```

```
[12]: #lets try to build a simple Ridge regression model. IESO data seems similar to
      ↪PJM data, so use IESO as training data and PJM as test
```

```
[13]: model = Ridge()
      df_ieso = df_ieso[df_ieso["P"] < 1000] # remove the two outliers
      X_train = df_ieso[['G','L']] # feature variables
      y_train = df_ieso[['P']] # target variables
      df_pjm_clean = df_pjm.dropna() # clean data by dropping NaN
      X_test = df_pjm_clean[['G','L']]
      y_test = df_pjm_clean[['P']]
      model.fit(X_train, y_train)
      predictions = model.predict(X_test)
      score = r2_score(y_test, predictions)
      print(score)
```

-404.72240535988794

```
[14]: # not a good model given bad R2 score. Let's plot the predicted and actual
      ↪ values to see what's going on
fig, ax = plt.subplots(figsize=(15, 5))
ax.plot(y_test, color='k', lw=3)
ax.plot(predictions, color='r', lw=2)
plt.show()
```



```
[15]: # Let's try using NYISO data as training and pjm as test
model = LinearRegression()
df_nyiso_clean = df_nyiso.dropna() # clean data by dropping NaN
X_train = df_nyiso_clean[['gas', 'load_rt']] # feature variables
y_train = df_nyiso_clean[['rt_lmp']] # target variable
X_test = df_pjm_clean[['G', 'L']]
y_test = df_pjm_clean[['P']]
model.fit(X_train, y_train)
predictions = model.predict(X_test)
score = r2_score(y_test, predictions)
print(score)
```

-466.45849946923306

```
[16]: #Still not working. let's try merging all the datasets and use 80% to train,
      ↪ 20% to test
df_nyiso_clean = df_nyiso2.dropna()
df_pjm_clean = df_pjm.dropna()
frames = [df_nyiso_clean, df_pjm_clean, df_ieso]
all_data = pd.concat(frames) # merge all the datasets
all_data = all_data[all_data["P"] < 1500] # drop outlier values
X = all_data[['G', 'L']]
y = all_data[['P']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
# build model
model = Ridge()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
score = r2_score(y_test, predictions)
score_ridge = model.score(X_test, y_test)
print(score)
print(score_ridge)
```

0.19053600155963524

0.19053600155963524

[17]: *# okay, more reasonable R2 score but still not good. Can we tune hyper-parameter*  
*↳ to improve model?*

```
scores = []
# loop around different values of alpha parameter
for alpha in [.01, .1, 1, 10, 1000, 10000]:
    model = Ridge(alpha=alpha)
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    scores.append(r2_score(y_test, predictions))
print(scores)
```

[0.19053612014988086, 0.19053610936913623, 0.19053600155963524,  
0.19053492325913612, 0.19041405230749597, 0.18912950779113458]

[18]: *# let's use 80% of pjm as training and 20% as test*

```
df_pjm_clean = df_pjm.dropna()
X = df_pjm_clean[['G', 'L']]
y = df_pjm_clean[['P']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
# build model
model = Ridge()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
score = r2_score(y_test, predictions)
score_ridge = model.score(X_test, y_test)
print(score)
print(score_ridge)
```

0.5601889156408217

0.5601889156408217

[23]: *# much better R-score. Let's use same strategy for IESO data*

```
X = df_ieso[['G', 'L']]
y = df_ieso[['P']]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
↳random_state=42)  
# build model  
model = Ridge()  
model.fit(X_train, y_train)  
predictions = model.predict(X_test)  
score = r2_score(y_test, predictions)  
score_ridge = model.score(X_test,y_test)  
print(score)  
print(score_ridge)
```

0.5677593610748015

0.5677593610748015

[ ]:

[ ]: